

Example: spec review of the brief 1 spec (v0.1)

EXAMPLE, for reference only. The spec review checklist applied to the example spec for brief 1, in the 25-minute clock, by the build team acting as reviewers. Question numbers match the checklist. The findings are real gaps in that spec. Your reviews will find different gaps; this shows how short and specific a good answer is, and what a BLOCKS finding looks like.

Clock as run: 7 min reading, 12 min answering (finished all ten groups with one minute spare), 4 min choosing, 2 min reading out.

A. Is the job real? (section 1)

#	Answer	Where
A1	YES: ticket text in, seven-field record plus valid out	1
A2	YES: 40 min median, 3 h after a weekend, 11% bounce-back	1
A3	YES: the agent, within five minutes, on the pre-filled form	1
A4	YES: the plugin and auto-routing are both things a reader assumes	1

B. Was the decision made, or assumed? (section 2)

#	Answer	Where
B1	YES: four totals, arithmetic in the matrix page	2
B2	YES: G1 attributed to the data owner by name and date; G5 answered "partly", which is honest but is not an answer	2
B3	YES: the DPA review completing is a this-year event	2
B4	YES: \$1,434 with 8 h × \$60 of ops labour; but the ops hours are the cost model's placeholder, not the team's number	7

C. Will the data owner sign this? (section 3)

#	Answer	Where
C1	YES: internal + personal data; D is allowed for it	3
C2	NO : pasted passwords: "OPEN: no mechanism yet". The intention is there, the mechanism is not	3
C3	YES: hash, output, 90 days, platform engineering	3, 8
C4	YES: log stores the hash, not the text. But the <i>output requested_action</i> is free text generated from the body; a pasted password could be echoed into it. Not addressed	3, 8
C5	YES: 600 / 400 / 80 / 20, five named checks	3

D. Can this actually be run? (section 4)

#	Answer	Where
D1	YES: 1B, r=16, 3 epochs; 3B ruled out with a reason	4
D2	CAN'T TELL : SKU, region and person are there, but the serving version is "a problem" and the fix (merge the adapter, serve on current Ollama or vLLM) is untested. As written, production runs on a pinned Ollama release that has since dropped the feature	4
D3	YES: <code>get_endpoint("tuned")</code> , swap by <code>.env</code>	4
D4	YES: 1,500-token cap, truncate the end, log it	4

#	Answer	Where
D5	YES: schema validated in the endpoint; <code>valid: false</code> and a blank form	4
D6	YES: 20 s, one retry, blank form, paging rule with a number	4

E. Retrieval and tools (section 5)

#	Answer	Where
E1–E3	n/a: none	5
E4	YES: ticket body named as untrusted, reaches the model as user content only, output is a validated record	5

F. Would you believe the quality number? (section 6)

#	Answer	Where
F1	YES: rubric named, per field	6
F2	YES: 20, stratified, lookalike-free, enforced by a test	6
F3	YES: four separate lines, plus per field	6
F4	YES: routing 16/20 today against 80% target; plausible	6
F5	NO, and the spec knows it: "too small for a pilot claim" is written down, and the 200-ticket shadow week is the answer. But the pilot target table already says "on 200 tickets" while every number in it is on 20	6, 9
F6	YES: 20 a week by the desk lead, misses become corrections	6

G. Does the money add up? (section 7)

#	Answer	Where
G1	YES: all seven	7
G2	YES: idle 99.7% acknowledged; the reason for hosting is gate G1, stated	2, 7
G3	YES: $p95 < 5$ s, measured on Day 2 S8	7
G4	YES: \$1,600 and "a decision, not autoscaling"	7

H. What happens when it is wrong? (section 8)

#	Answer	Where
H1	YES: pre-filled, never submitted; urgency always the agent's	8
H2	YES: twelve fields including adapter fingerprint	8
H3	YES: webhook off, 30 seconds, two named people	8
H4	YES: re-register the previous adapter, under five minutes, tested on Day 2	8

I. Will it survive us leaving? (section 9)

#	Answer	Where
I1	NO: "in principle", manager has not agreed, no GPU experience. This is gate G5 answered "partly" and never resolved	0, 2, 9
I2	YES: MRB desk, four agents, six weeks, four measures	9
I3	YES: data owner, security, desk lead, the owner's manager	9

J. The manager's three questions

#	One-line answer from the spec
J1	\$1,434 a month, almost all of it the VM and the engineer's time; saves agent minutes we have not yet measured (the shadow week measures them). Honest: at this volume a vendor would cost \$5 a month, and the DPA review is what stands between us and that
J2	A wrong queue or urgency pre-filled on the form: the agent sees it before it routes. Worst case we can name: a pasted password echoed into <code>requested_action</code> and stored in the log
J3	Faisal, in principle. Nobody has agreed that in writing

Findings block

#	Severity	Finding	Question
1	BLOCKS	There is no owner. "In principle" is not a name a manager will accept, and gate G5 was answered "partly" and carried through nine sections unresolved. Until someone whose job includes this VM says yes, this is a lab, not a deployment	I1, B2
2	BLOCKS	The must-never-enter mechanism does not exist. The spec's own reason for choosing D over B is that ticket text contains pasted passwords; it then sends that text to the model unredacted and can echo it into a stored output field. The data owner will not sign section 3 as written	C2, C4
3	BEFORE PILOT	The serving stack is not runnable as specified: the adapter needs an Ollama release that has since dropped adapter support. The merged-weights path is the right fix and it is untested. It must be tested before the pilot, not "on Day 5 S27" if Day 5 runs late	D2

Also noted, not in the top three: F5: the pilot target table claims 200-ticket targets on 20-ticket evidence; write "target on the 200-ticket shadow week" as the column heading. B4: the ops hours are the cost model's placeholder; ask Faisal for his.

The best thing about this spec: section 8. A wrong answer cannot reach a queue without an agent seeing it, urgency is explicitly the agent's field because the eval says the model cannot be trusted with it, and the model being down leaves the desk exactly where it is today. That is what "deployment-ready" is supposed to mean, and it should not be lost in revision.