

Capstone quick card

A ticket goes in. It is looked up. The model answers. The answer is checked. A decision is made: propose, refuse, or ask a person. Everything is logged. **You change one file:** `capstone/my_usecase.py`. Nothing in the scaffold writes anything, so your capstone is a graph you draw and run.

Time	Check	Commands / what to show
10:30	Pick a brief	1 ticket to structured record, 2 HSE procedure assistant, 3 nameplate and diagram capture, 4 ticket to work order, 5 similar-ticket assist. Same brief twice is fine.
10:45	1 Run it untouched	<code>python -m capstone.run status</code> <code>python -m capstone.run ticket INC-005310</code> Read every row. A FALLBACK row is fine; say why.
11:15	2 Your edit changes the answer	Edit <code>my_usecase.py</code> . Run the ticket command again. Point at your brief in the output.
11:40	3 One wrong answer caught, one number	<code>check_output</code> fires on a bad answer. <code>python -m capstone.run eval</code> Write the number down, with what it does NOT show (n = 20).

The five parts of the file

Part	In plain words
SETTINGS	Name of your use case and its limits
<code>gather_context</code>	What to look up before asking the model
<code>build_messages</code>	What to ask the model
<code>check_output</code>	Is this answer acceptable? If not, what is wrong?
<code>decide</code>	Propose it, refuse it, or ask a person. Nothing writes.

Building on your own?

Copy the file once so you do not overwrite each other: `cp capstone/my_usecase.py capstone/my_usecase_yourname.py`, then add `--usecase capstone.my_usecase_yourname` to each command.

Stuck? 1) Copy `capstone/examples/brief5_similar_tickets.py` and change the names (78 lines). 2) Read the saved output in `facilitator/prebaked_outputs/capstone/`. 3) Ask.

What you show at 14:00: a working run, your graph picture (nodes, arrows, who decides, caps, the gate), your checklist verdict, templates 1 and 4, and one wrong answer being caught.